

# **Concurrency and Parallelism**

## **CSE 3341**

**Reading:**  
**Chapter 13 in 4th Edition**

# Outline

- **Motivation & examples**
  - Threads, shared memory, & synchronization
    - How do locks work?
  - Data races (a lower-level property)
  - How do data race detectors work?
  - Atomicity (a higher-level property)
  - Concurrency exceptions & Summary
- Extra:
- Double-checked locking

# Motivation

- A program is **concurrent** if it may have more than one active execution context – “threads”
  - A **thread** is the smallest sequence of programmed instructions that can be managed independently by the OS scheduler
- We do this to
  - Capture the logical structure of a program
  - Exploit parallelism for speed
  - Cope with physical distribution

# Concurrent vs Parallel vs Distributed

- A concurrent system is **parallel** if more than one task can be physically active at once (this requires  $>1$  processor)
- A parallel system is **distributed** if processors/devices are physically separated in the real world
- Parallelism vs switching is a performance and implementation issue, not a semantic one

**Hardware becoming more parallel**  
**(*more* instead of *faster* cores)**

**Software must become more parallel**

“From my perspective, parallelism is the biggest challenge since high-level programming languages. It’s the biggest thing in 50 years because industry is betting its future that parallel programming will be useful.

“Industry is building parallel hardware, assuming people can use it. And I think there's a chance they'll fail since **the software is not necessarily in place**. So this is a **gigantic challenge facing the computer science community**. If we miss this opportunity, it's going to be bad for the industry.”

—David Patterson, ACM Queue interview, 2006

# Shared-memory programming

## ■ Imperative programs

- Java, C#, C, C++, Python, Ruby

## ■ Threads

- Shared, mutable state
- Synchronization primitives:
  - Lock acquire & release
  - Monitor wait & notify
  - Thread start & join

# What are the possible behaviors?

```
int x = 1;
```

**T1:**

```
t = x;  
t = t + 1;  
x = t;
```

**T2:**

```
t = x;  
t = t + 1;  
x = t;
```



# Atomicity & determinism

```
int x = 1;
```

**T1:**

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

**T2:**

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

# What are the possible behaviors?

```
int x = 1;
```

**T1:**

```
t = x;  
t = t + 1;  
x = t;
```

**T2:**

```
t = x;  
t = t * 2;  
x = t;
```

# Atomicity (still nondeterminism)

```
int x = 1;
```

T1:

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

T2:

```
synchronized (m) {  
    t = x;  
    t = t * 2;  
    x = t;  
}
```

# What are the possible behaviors?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = 42;  
flag = true;
```

**T2:**

```
if (flag) {  
    int t = data;  
    print(t);  
}
```

# Outline

- Motivation & examples
  - **Threads, shared memory, & synchronization**
    - **How do locks work?**
  - Data races (a lower-level property)
  - How do data race detectors work?
  - Atomicity (a higher-level property)
  - Concurrency exceptions & Summary
- Extra:
- Double-checked locking

# Threads, shared memory, & locks

Each thread:

- has its own stack
- shares memory with other threads in same process (same virtual address space)

Compiler compiles code as though it were single-threaded!

Synchronization operations (e.g., lock acquire and release) order accesses to shared memory, providing:

- mutual exclusion
- ordering and visibility

# Locks in Java

```
int x = 1;
```

T1:

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

T2:

```
synchronized (m) {  
    t = x;  
    t = t * 2;  
    x = t;  
}
```

# Locks in Java

```
int x = 1;
```

**T1:**

```
acquire(m) ;  
    t = x;  
    t = t + 1;  
    x = t;  
release(m) ;
```

**T2:**

```
acquire(m) ;  
    t = x;  
    t = t * 2;  
    x = t;  
release(m) ;
```




# Locks in Java

```
int x = 1;
```

T1:

```
acquire(m.lockBit) ;  
  t = x;  
  t = t + 1;  
  x = t;  
release(m.lockBit) ;
```



T2:

```
acquire(m.lockBit) ;  
  t = x;  
  t = t * 2;  
  x = t;  
release(m.lockBit) ;
```

# Locks in Java

```
int x = 1;
```

T1:

```
while (m.lockBit != 0) {}  
m.lockBit = 1;  
    t = x;  
    t = t + 1;  
    x = t;  
m.lockBit = 0;
```

T2:

```
while (m.lockBit != 0) {}  
m.lockBit = 1;  
    t = x;  
    t = t * 2;  
    x = t;  
m.lockBit = 0;
```

Possible implementation of locks?

# Locks in Java

```
int x = 1;
```

T1:

```
while (!TAS(m.lockBit,0,1)) {}  
    t = x;  
    t = t + 1;  
    x = t;  
m.lockBit = 0;
```

T2:

```
while (!TAS(m.lockBit,0,1)) {}  
    t = x;  
    t = t * 2;  
    x = t;  
m.lockBit = 0;
```

Need an atomic operation like **test-and-set (TAS)**

# Locks in Java

```
int x = 1;
```

**T1:**

```
while (!TAS(m.lockBit,0,1)) {}  
    t = x;  
    t = t + 1;  
    x = t;  
memory_fence;  
m.lockBit = 0;
```

**T2:**

```
while (!TAS(m.lockBit,0,1)) {}  
    t = x;  
    t = t * 2;  
    x = t;  
memory_fence;  
m.lockBit = 0;
```

- Fence needed for visibility (related to happens-before relationship, discussed later)
- Also: compiler obeys “roach motel” rules: can move operations into but not out of atomic blocks

# Locks in Java

```
int x = 1;
```

T1:

```
while (!TAS(m.lockBit,0,1)) {}  
    t = x;  
    t = t + 1;  
    x = t;  
memory_fence;  
m.lockBit = 0;
```

T2:


```
while (!TAS(m.lockBit,0,1)) {}  
    t = x;  
    t = t * 2;  
    x = t;  
memory_fence;  
m.lockBit = 0;
```

- Java locks are **reentrant**, so more than one bit is actually used (to keep track of nesting depth)
- Also, **spin (non-blocking) locks** are converted to **blocking locks** if there's contention


# Outline

- Motivation & examples
  - Threads, shared memory, & synchronization
    - How do locks work?
  - **Data races (a lower-level property)**
  - How do data race detectors work?
  - Atomicity (a higher-level property)
  - Concurrency exceptions & Summary
- Extra:
- Double-checked locking

# Northeast Blackout of 2003



# Northeast Blackout of 2003



50 million  
people



# Northeast Blackout of 2003

- Energy Management System
  - Alarm and Event Processing Routine (1 MLOC)



# Northeast Blackout of 2003

## ■ Energy Management System

- Alarm and Event Processing Routine (1 MLOC)

## ■ Post-mortem analysis: 8 weeks

"This fault was so deeply embedded, it took them weeks of poring through millions of lines of code and data to find it." –Ralph DiNicola, FirstEnergy

ISAT GeoStar 45

# Northeast Blackout of 2003

- Race condition
  - Two threads writing to data structure simultaneously
- Usually occurs without error
  - Small window for causing data corruption

ISAT GeoStar 45

# What is a data race?

- Two accesses to same variable
- At least one is a write

Not well-synchronized

(not ordered by *happens-before* relationship)

Or: accesses can happen simultaneously

# DRF0-based memory models

Modern language **memory models** (via compiler+hardware) guarantee the following relationship:

Data race freedom  $\rightarrow$  Sequential consistency

However: Data race  $\rightarrow$  Weak or undefined semantics!

Sequential consistency: instructions appear to execute in an order that respects program order

# Is there a data race?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = 42;  
flag = true;
```

**T2:**

```
if (flag)  
    t = data;
```

# Is there a data race?

Initially:

```
int data = 0;
```

```
boolean flag = false;
```

**T1:**

```
data = 42;
```

```
flag = true;
```

**T2:**

```
if (flag)
```

```
    t = data;
```



# Possible behavior

Initially:

```
int data = 0;
```

```
boolean flag = false;
```

**T1:**

```
flag = true;
```

```
data = 42;
```

**T2:**

```
if (flag)  
    t = data;
```



# Possible behavior

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = 42;  
flag = true;
```

**T2:**

```
t2 = data;
```

```
if (flag)  
    t = t2;
```

# Is there a data race?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1**:

```
data = ...;  
synchronized (m) {  
    flag = true;  
}
```

**T2**:

```
boolean f;  
synchronized (m) {  
    f = flag;  
}  
if (f)  
    ... = data;
```

# Is there a data race?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = ...;  
acquire(m) ;  
    flag = true;  
release(m) ;
```

**T2:**

```
boolean f;  
acquire(m) ;  
    f = flag;  
release(m) ;  
if (f)  
    ... = data;
```

*Happens-before  
relationship*



# Is there a data race?

Initially:

```
int data = 0;
```

```
volatile boolean flag = false;
```


**T1:**

```
data = ...;  
flag = true;
```

**T2:**

```
if (flag)  
    ... = data;
```

*Happens-before  
relationship*



# What can this program print?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = 42;  
flag = true;
```

**T2:**

```
while (!flag) { }  
print(data);
```

# What can this program print?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = 42;  
flag = true;
```

**T2:**

```
boolean f;  
do {  
    f = flag;  
} while (!f);  
int d = data;  
print(d);
```

# What can this program print?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
flag = true;
```

```
data = 42;
```

**T2:**

```
boolean f;  
do {  
    f = flag;  
} while (!f);  
int d = data;  
print(d);
```

# What can this program print?

Initially:

```
int data = 0;  
boolean flag = false;
```

**T1:**

```
data = 42;  
flag = true;
```

**T2:**

```
int d = data;  
  
boolean f;  
do {  
    f = flag;  
} while (!f);  
print(d);
```



# Outline

- Motivation & examples
  - Threads, shared memory, & synchronization
    - How do locks work?
  - Data races (a lower-level property)
  - **How do data race detectors work?**
  - Atomicity (a higher-level property)
  - Concurrency exceptions & Summary
- Extra:
- Double-checked locking

# Data Races

- Two accesses to same variable (one is a write)
- One access doesn't **happen before** the other
  - Program order
  - Synchronization order
    - Acquire-release
    - Wait-notify
    - Fork-join
    - Volatile read-write

# Data Races

- Two accesses to same variable (one is a write)
- One access doesn't **happen before** the other
  - Program order
  - Synchronization order
    - Acquire-release
    - Wait-notify
    - Fork-join
    - Volatile read-write

Thread A

write x



unlock m

Thread B

# Data Races

- Two accesses to same variable (one is a write)
- One access doesn't **happen before** the other
  - Program order
  - **Synchronization order**
    - Acquire-release
    - Wait-notify
    - Fork-join
    - Volatile read-write

Thread A

write x

unlock m

Thread B

lock m

write x



# Data Races

- Two accesses to same variable (one is a write)
- One access doesn't **happen before** the other
  - Program order
  - Synchronization order
    - Acquire-release
    - Wait-notify
    - Fork-join
    - Volatile read-write

## Thread A

write x



unlock m



read x

## Thread B

lock m




write x



# Data Races

- Two accesses to same variable (one is a write)
- One access doesn't **happen before** the other
  - Program order
  - Synchronization order
    - Acquire-release
    - Wait-notify
    - Fork-join
    - Volatile read-write




# Data Race Detection: Possible Strategies

How do dynamic data race detectors work?

- Vector clocks: check happens-before
- Collision analysis: simultaneous conflicting accesses
- Lockset: assume and check locking discipline

# What is a Vector Clock?

- Each thread  $T$  maintains its own logical clock 'c'
  - Initially,  $c = 0$  when  $T$  starts
  - Incremented at synchronization release operations
    - unlock  $m$ , volatile write
- Vector clock is a vector of logical clocks





# Vector Clock and Happens Before

$$V_1 \sqsubseteq V_2 \quad \text{iff} \quad \forall t \quad V_1(t) \leq V_2(t)$$

A	B	C
1	2	3

$\sqsubseteq$

A	B	C
5	6	7

true

A	B	C
4	5	3

$\sqsubseteq$

A	B	C
4	5	3

true

A	B	C
4	5	3

$\sqsubseteq$

A	B	C
4	2	3

false

# Vector Clock-Based Race Detection

## Note:

Starting the Vector Clock algorithm, I would initialize the clocks to  $[1, 0]$  and  $[0, 1]$ .


The intuition here is that Thread A knows it has started execution, but does not know that thread B has started execution.

When a thread acquires the lock, information is gained from the last thread that held the lock. On the first lock, no information is gained.


The next few slides are meant to represent the state after several locks/unlocks, so the vector clocks have advanced significantly.

# Vector Clock-Based Race Detection

Thread A




Thread B



Vector clocks


# Vector Clock-Based Race Detection

Thread A



Thread A's logical  
time

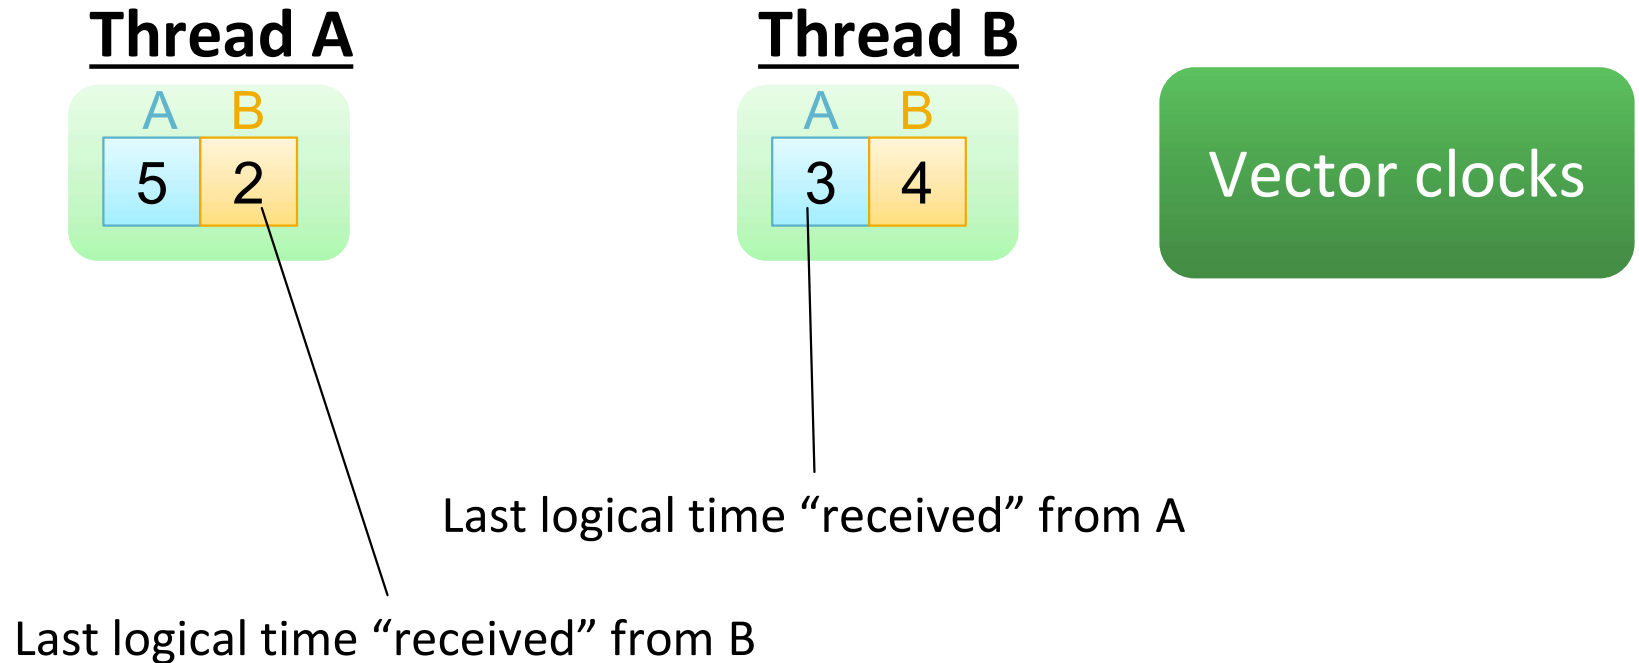
Thread B



Thread B's logical time


Vector clocks

# Vector Clock-Based Race Detection




# Vector Clock-Based Race Detection

Thread A



Thread B




5	2
---	---



 write x

# Vector Clock-Based Race Detection

Thread A





Thread B



write x


unlock m





Increment  
clock

# Vector Clock-Based Race Detection

Thread A




Thread B



write x

unlock m




Increment  
clock

lock m





# Vector Clock-Based Race Detection

Thread A




Thread B




write x

unlock m



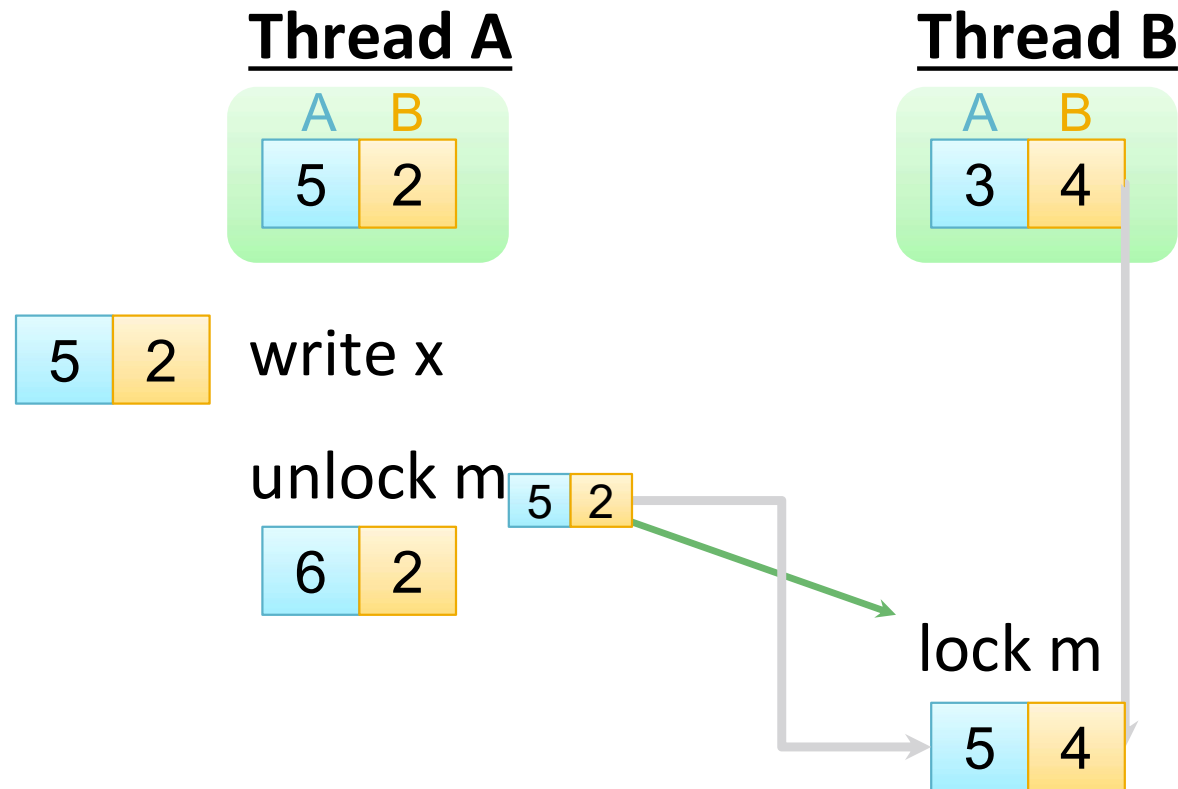
lock m



Increment  
clock

Join  
clocks

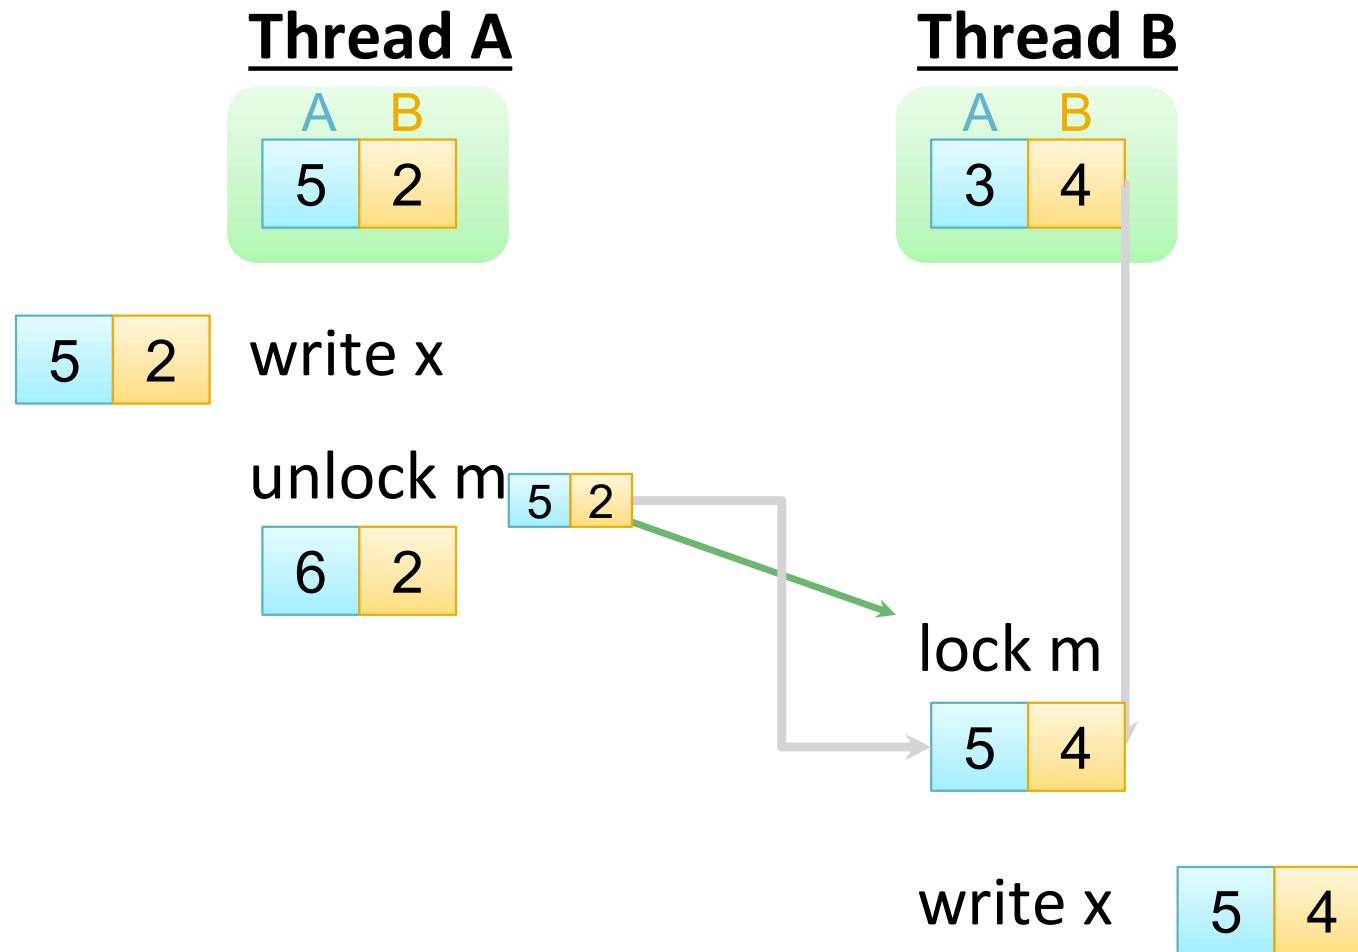
# Vector Clock-Based Race Detection



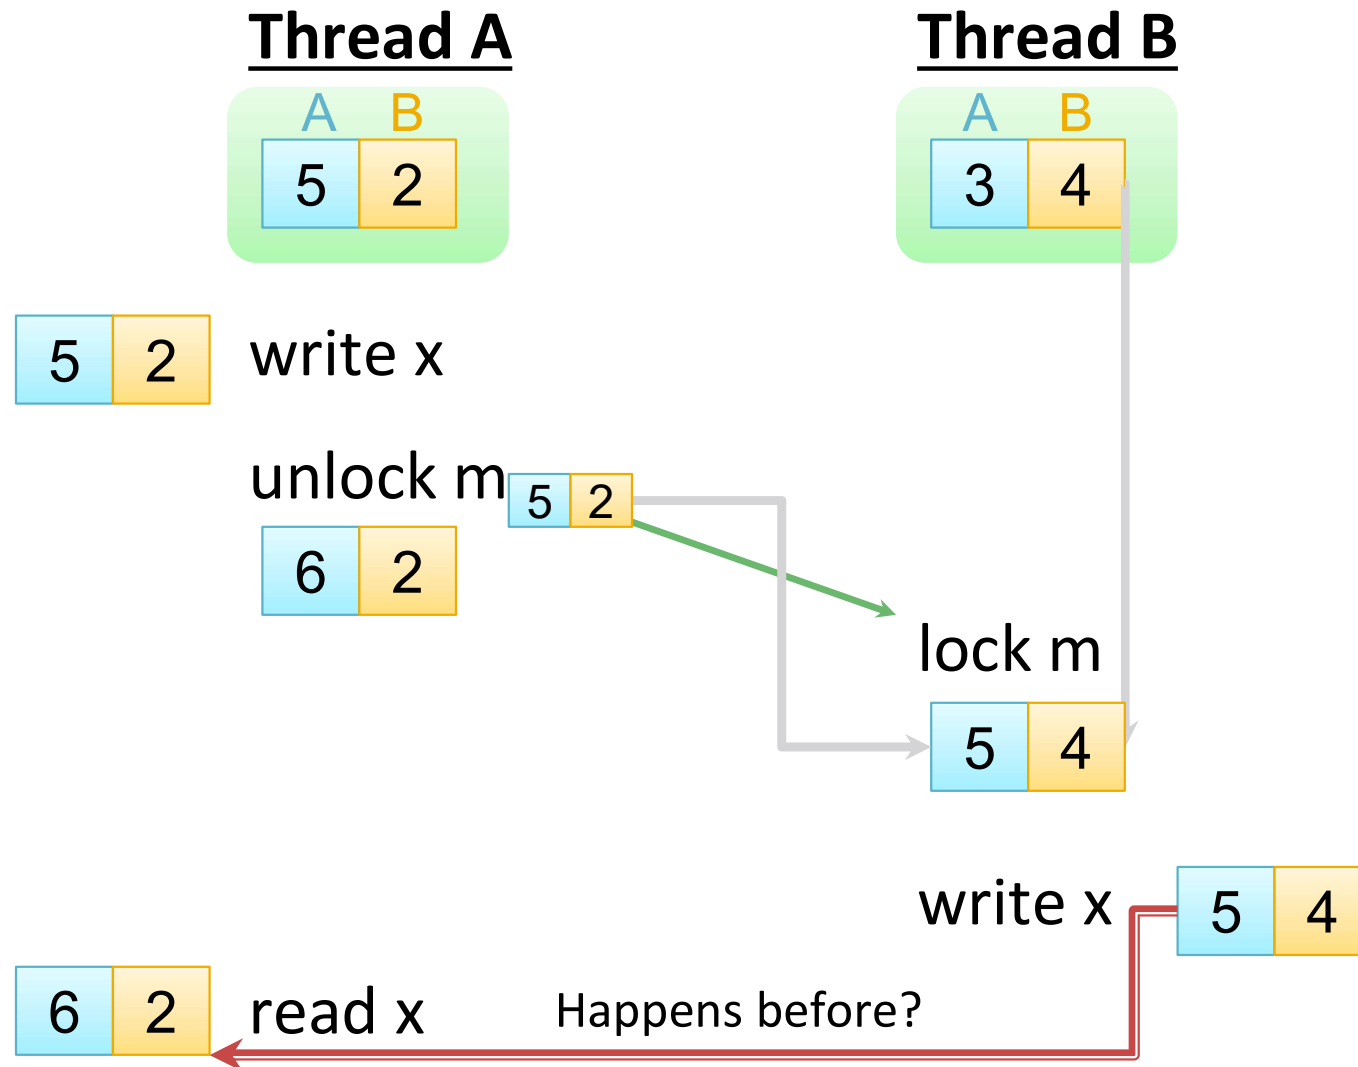
$O(n)$  time

$n = \#$  of threads

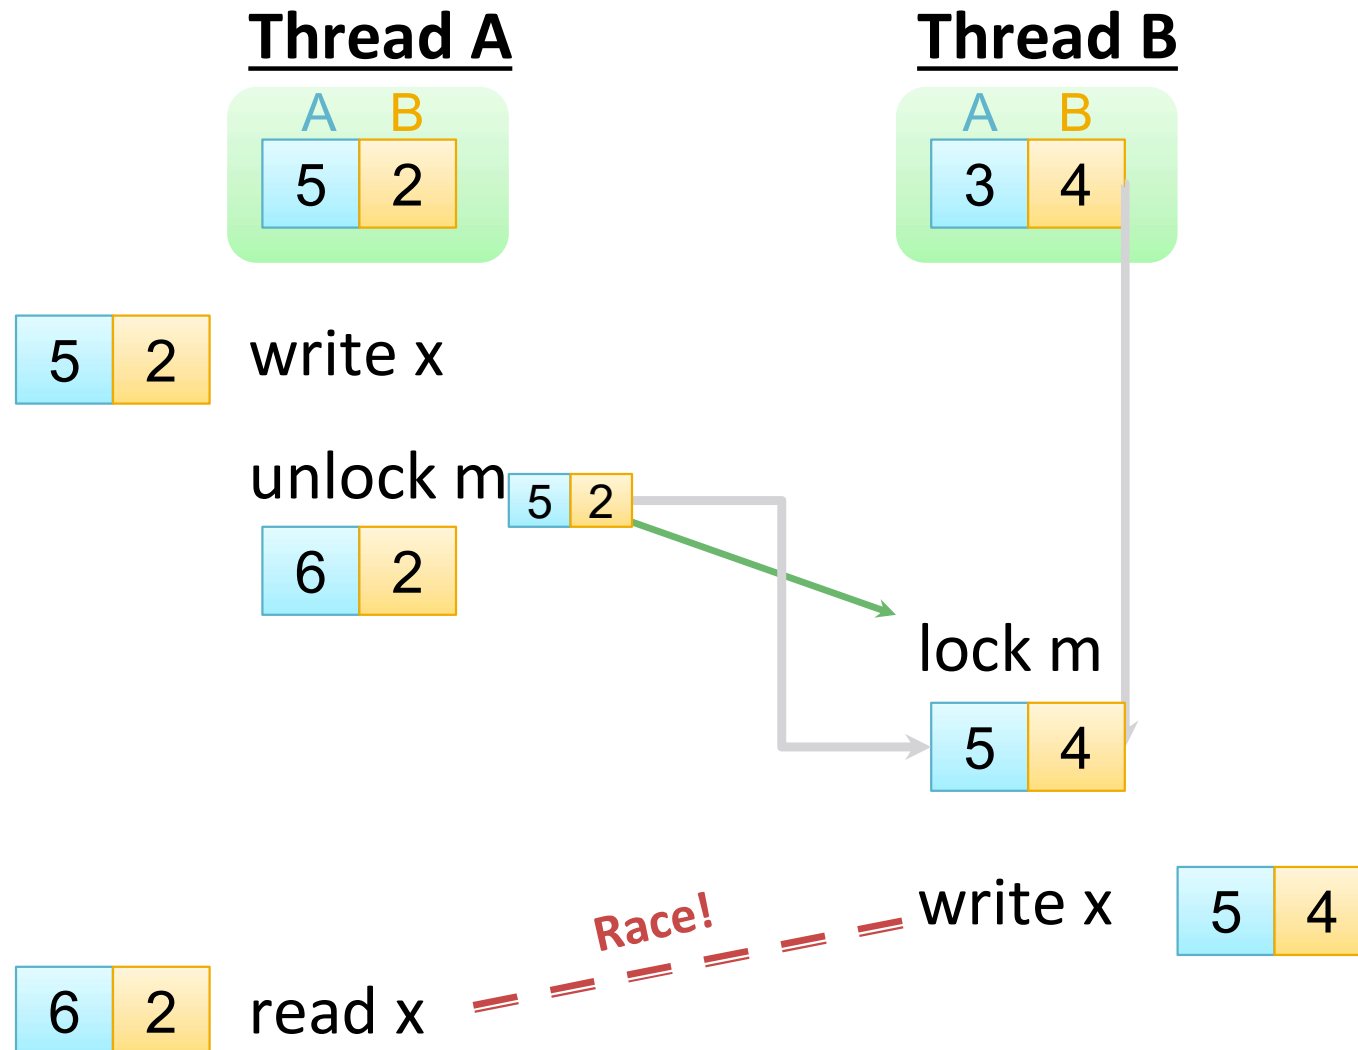
# Vector Clock-Based Race Detection



# Vector Clock-Based Race Detection



# Vector Clock-Based Race Detection



# Soundness and Precision

## ■ Soundness

- No false negatives, i.e., all data races in **that** execution will be reported.

## ■ Precision

- No false positives, i.e., all reports are true data races.

---

These are not standard terms across all domains, e.g., architects might refer to these properties as complete and sound.

# Vector Clock-Based Race Detection

- Tracks happens-before: sound & precise
  - **80X** slowdown
  - Each analysis step:  **$O(n)$**  time (n = # of threads)
- FastTrack [Flanagan & Freund '09]
  - Reads & writes (97%):  **$O(1)$**  time
  - Synchronization (3%):  **$O(n)$**  time
  - **8X** slowdown

# Vector Clock-Based Race Detection

- Tracks happens-before: sound & precise
  - **80X** slowdown
  - Each analysis step:  **$O(n)$**  time (n = # of threads)

- FastTrack [Flanagan & Freund '09]
  - Reads & writes (97%):  **$O(1)$**  time
  - Synchronization (3%):  **$O(n)$**  time
  - **8X** slowdown




# FastTrack's Insight

- In a data-race-free (DRF) program, writes are totally ordered
  - Can maintain only the “last” writer

# Vector Clock-Based Race Detection

Thread A



write x

unlock m


Thread B



lock m

write x

read x



# Vector Clock-Based Race Detection

Thread A

A	B
5	2

write x 5@A

unlock m


read x

Thread B

A	B
3	4

lock m

write x



# Vector Clock-Based Race Detection

Thread A

A	B
5	2

write x 5@A

unlock m


read x

Thread B

A	B
3	4

lock m

write x



# Vector Clock-Based Race Detection

Thread A

A	B
5	2

write x 5@A

unlock m

6	2
---	---

5	2
---	---


Thread B

A	B
3	4

lock m

write x

read x



# Vector Clock-Based Race Detection

Thread A

A	B
5	2

write x

5@A

unlock m

6	2
---	---

5	2
---	---


Thread B

A	B
3	4

lock m

write x

read x



# Vector Clock-Based Race Detection

## Thread A

A	B
5	2

write x 5@A

unlock m 5 2

6	2
---	---

## Thread B

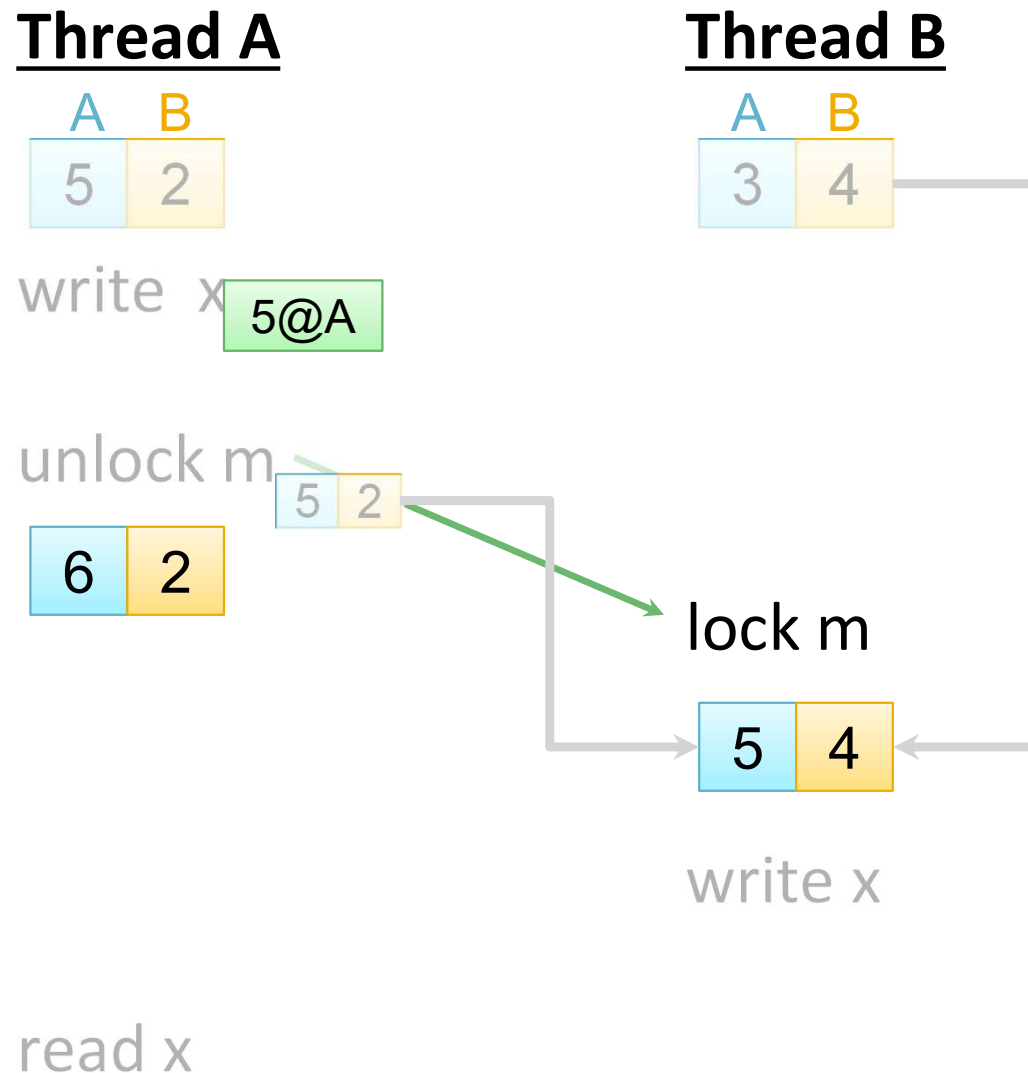
A	B
3	4

lock m

5	4
---	---

write x

read x



# Vector Clock-Based Race Detection

Thread A

A	B
5	2

write x

5@A

unlock m

6	2
---	---

5	2
---	---

Thread B

A	B
3	4

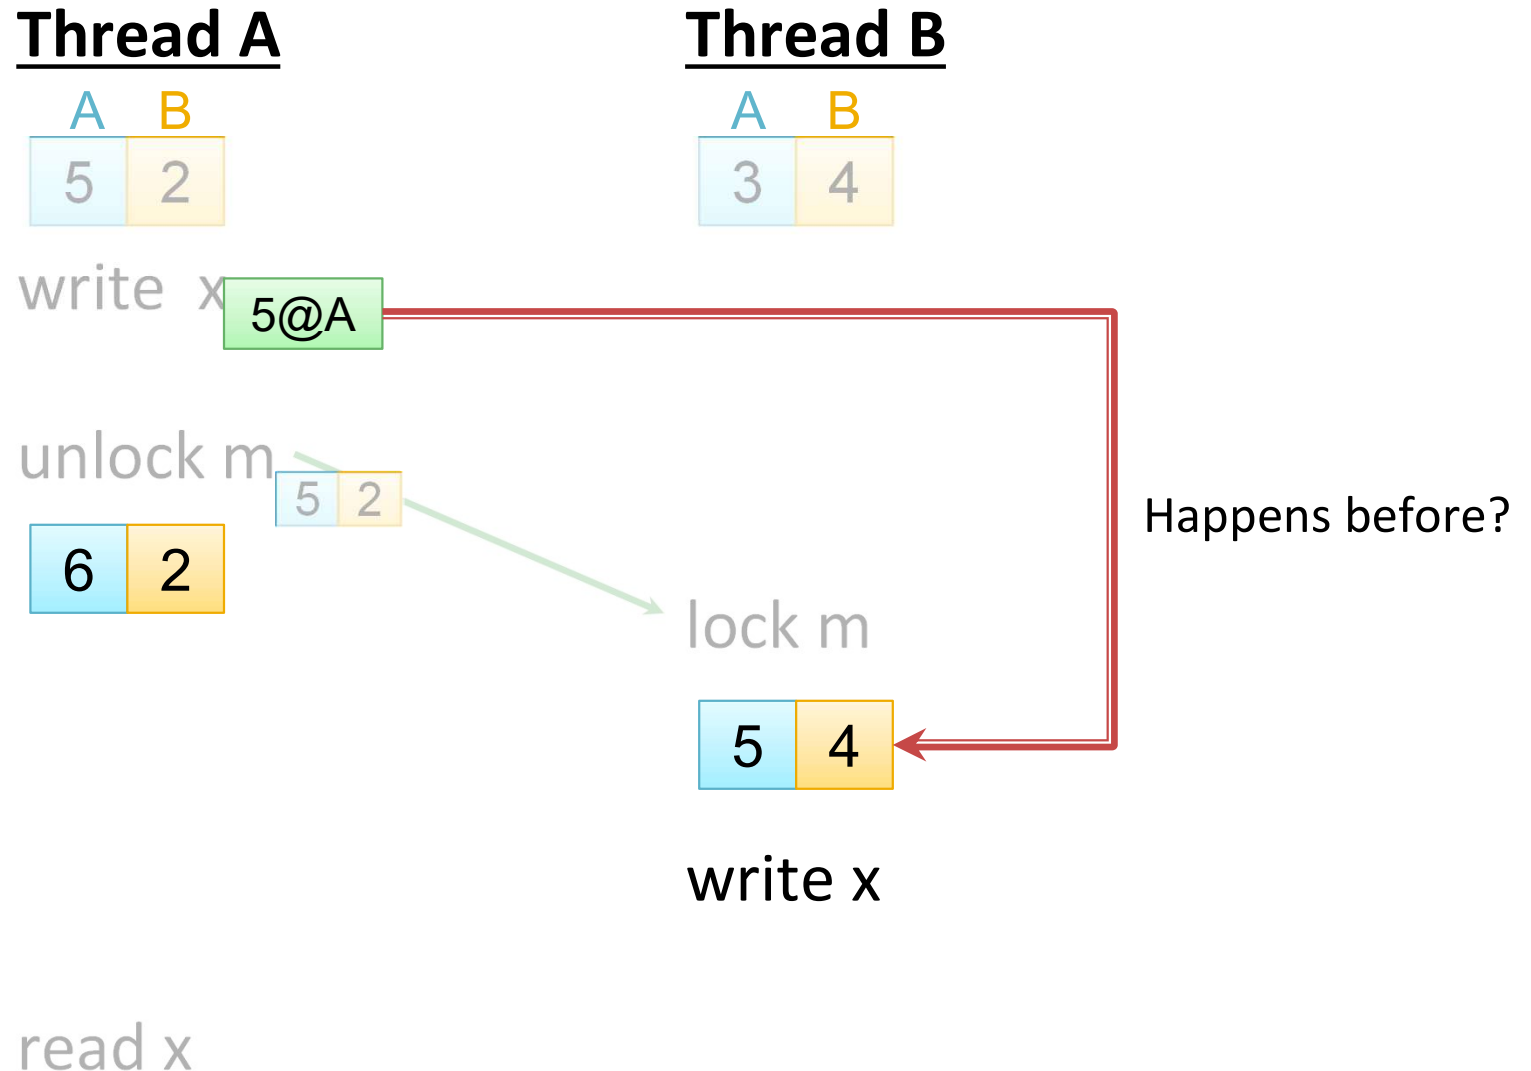
lock m

5	4
---	---

write x

read x

Happens before?





# Vector Clock-Based Race Detection

## Thread A

A	B
5	2

write x 5@A

unlock m

6	2
---	---

read x

## Thread B

A	B
3	4

lock m

5	4
---	---

write x 4@B

5	2
---	---



# Vector Clock-Based Race Detection

## Thread A

A	B
5	2

write x 5@A

unlock m 5 2

6	2
---	---

read x

## Thread B


A	B
3	4

lock m

5	4
---	---

write x 4@B

Happens before?



# Vector Clock-Based Race Detection

## Thread A

A	B
5	2

write x 5@A

unlock m 5 2

6	2
---	---

read x

## Thread B

A	B
3	4

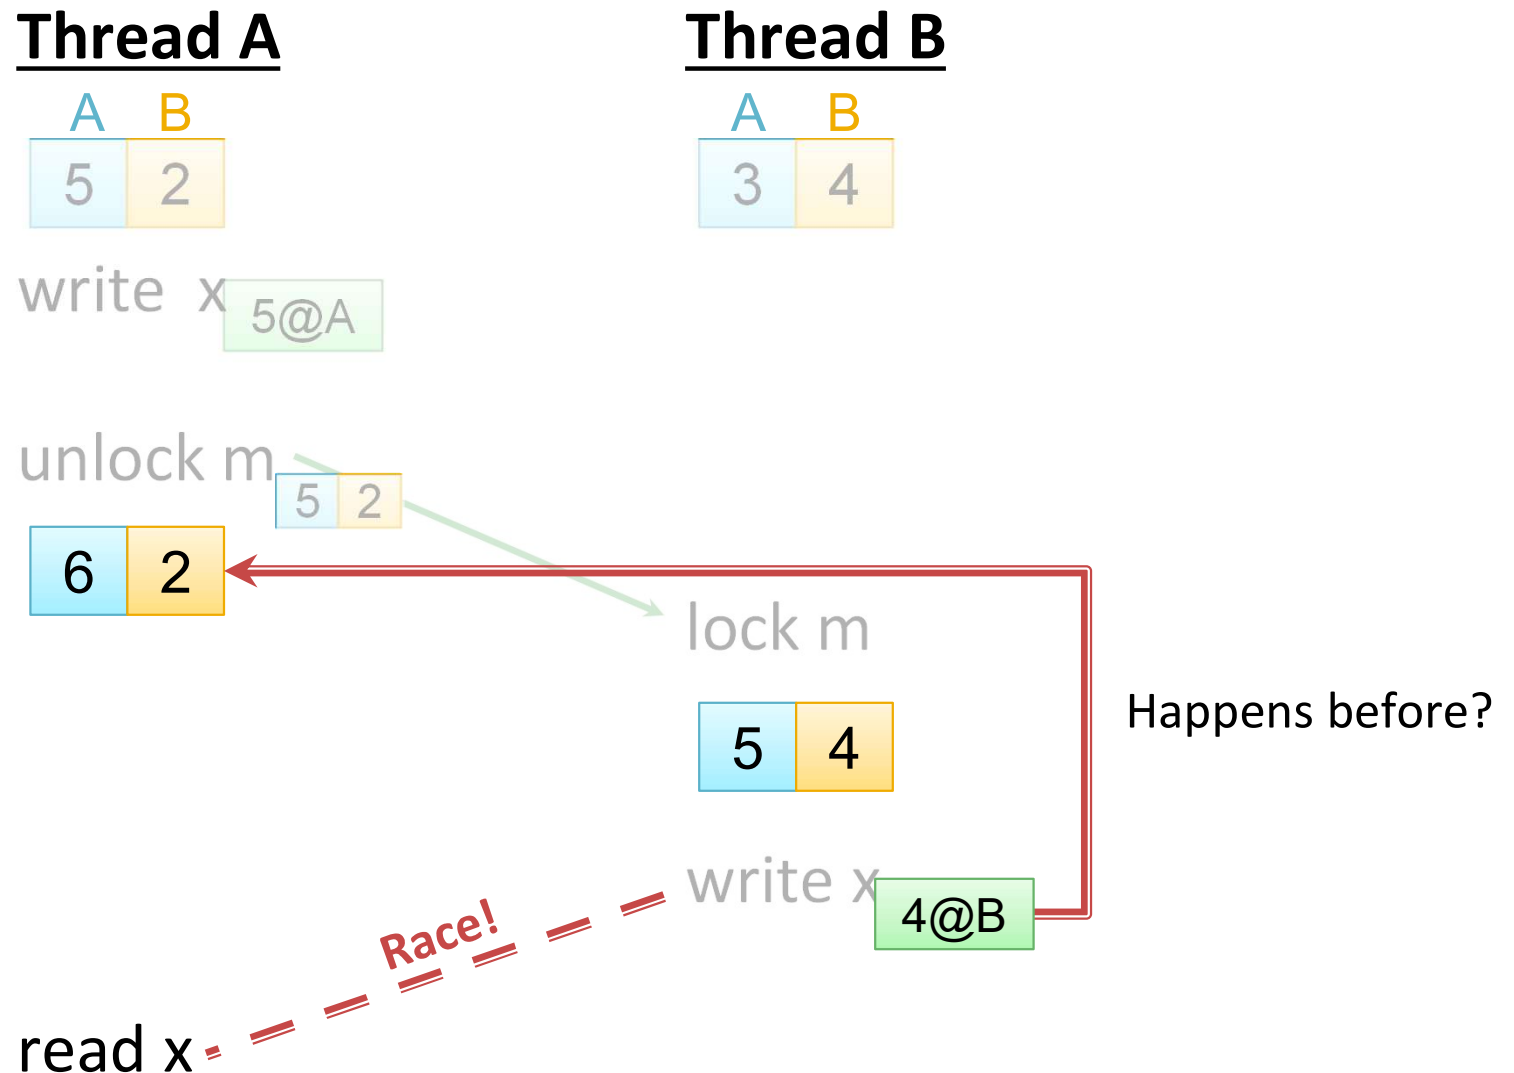
lock m

5	4
---	---

write x 4@B

Happens before?

**Race!**



# Collision analysis

Basic idea: Make two conflicting accesses happen at the same time

- (1) Pause one thread at memory access to  $x$
- (2) Catch other threads that access  $x$  (in a conflicting way) in the meantime






Can be automatic (with instrumentation or hardware) or manual (programmer adds code)

# What is a Lockset?

- Keeps track of the locks associated with each thread and program variable

# What is a Lockset?

- Keeps track of the locks associated with each thread and program variable

<u>Thread A</u>	<u>Lockset<sub>A</sub></u>
lock m	 $L = \{ \}$
write x	 $L = \{m\}$
lock n	
write y	 $L = \{m, n\}$
unlock n	
unlock m	 $L = \{m\}$
read x	 $L = \{ \}$

# Lockset Algorithms

- Two accesses from different threads with non-intersecting locksets form a data race
- Can detect **more data races** than vector clocks
  - A property referred to as **coverage**

# Is there a data race on variable $y$ ?

## Observed interleaving

### Thread A

$y = y + 1$   
lock  $m$   
 $v = v + 1$   
unlock  $m$

### Thread B

lock  $m$   
 $v = v + 1$   
unlock  $m$   
 $y = y + 1$



Happens before



# Is there a data race on variable $y$ ?

## Observed interleaving

### Thread A

$y = y + 1$   
lock  $m$   
 $v = v + 1$   
unlock  $m$

### Thread B

lock  $m$   
 $v = v + 1$   
unlock  $m$   
 $y = y + 1$

Happens before?



Data race with  
vector clock

Data race with  
lockset

# Alternate Interleaving

## Thread A

$y = y + 1$

lock m

$v = v + 1$

unlock m

*Happens before*



No data race with  
vector clock

## Thread B

lock m

$v = v + 1$

unlock m

$y = y + 1$

# Alternate Interleaving

## Thread A

$y = y + 1$

lock m

$v = v + 1$

unlock m

*Happens before*



Data race with  
lockset

## Thread B

lock m

$v = v + 1$

unlock m

$y = y + 1$

# Data Race on $y$ with Lockset

## Thread A

$y = y + 1$

lock  $m$

$v = v + 1$

unlock  $m$

## Thread B

lock  $m$

$v = v + 1$

unlock  $m$

$y = y + 1$

## Lockset

$L_y = \{ \}$

$L_v = \{m\}$

$L_v = \{m\}$

$L_y = \{ \}$

# Lockset Algorithms are Imprecise

```
volatile boolean flag = false;  
Object x;
```

## Thread A

```
x = new Object();  
flag = true;
```

## Thread B

```
while (!flag);  
x.print();
```

*Happens before*



# Lockset Algorithms are Imprecise


volatile boolean flag = false;  
Object x;

## Thread A

x = new Object();  
flag = true;

## Thread B

while (!flag);



B is  
blocked

while (!flag);  
x.print();

*Happens before*



# Vector Clock vs Lockset

- Lockset algorithms
  - Generally unsound and imprecise
  - Better coverage
- Vector clock algorithms
  - Sound and precise
  - Limited coverage

# Outline

- Motivation & examples
  - Threads, shared memory, & synchronization
    - How do locks work?
  - Data races (a lower-level property)
  - How do data race detectors work?
  - **Atomicity (a higher-level property)**
  - Concurrency exceptions & Summary
- Extra:
- Double-checked locking



# Atomicity

Operations appear to happen all at once or not at all

Serializability – execution equivalent to some serial execution of atomic blocks

# Atomicity violation?

```
int x = 1;
```

**T1:**

```
t = x;  
t = t + 1;  
x = t;
```

**T2:**

```
t = x;  
t = t * 2;  
x = t;
```

# Atomicity violation?

```
int x = 1;
```

**T1:**

```
t = x;
```

```
t = t + 1;
```

```
x = t;
```

**T2:**

```
t = x;
```

```
t = t * 2;
```

```
x = t;
```

# Atomicity violation?

```
int x = 1;
```

T1:

```
synchronized(m) {  
    t = x;  
}
```

```
t = t + 1;
```

```
synchronized(m) {  
    x = t;  
}
```

T2:

```
synchronized(m) {  
    t = x;  
}
```

```
t = t * 2;
```

```
synchronized(m) {  
    x = t;  
}
```

Still an atomicity violation

# Atomicity

```
int x = 1;
```

**T1:**

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

**T2:**

```
synchronized (m) {  
    t = x;  
    t = t * 2;  
    x = t;  
}
```

# Atomicity

```
int x = 1;
```

**T1:**

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

**T2:**

```
synchronized (m) {  
    t = x;  
    t = t * 2;  
    x = t;  
}
```

# Atomicity with different outcome

```
int x = 1;
```

**T1:**

```
synchronized (m) {  
    t = x;  
    t = t + 1;  
    x = t;  
}
```

**T2:**

```
synchronized (m) {  
    t = x;  
    t = t * 2;  
    x = t;  
}
```

# Atomicity violation without data races

```
class Vector {  
    synchronized boolean contains(Object o) { ... }  
    synchronized void add(Object o) { ... }  
}
```



# Atomicity violation without data races

```
class Vector {  
    synchronized boolean contains(Object o) { ... }  
    synchronized void add(Object o) { ... }  
}
```

```
class Set {  
    Vector vector;  
    void add(Object o) {  
        if (!vector.contains(o)) {  
            vector.add(o);  
        }  
    }  
}
```

# Atomicity violation without data races

```
class Vector {  
    synchronized boolean contains(Object o) { ... }  
    synchronized void add(Object o) { ... }  
}
```

```
class Set {  
    Vector vector;  
    synchronized void add(Object o) {  
        if (!vector.contains(o)) {  
            vector.add(o);  
        }  
    }  
}
```

# Atomicity violation without data races

```
class Vector {  
    synchronized boolean contains(Object o) { ... }  
    synchronized void add(Object o) { ... }  
}
```

```
class Set {  
    Vector vector;  
    void add(Object o) {  
        atomic {  
            if (!vector.contains(o)) {  
                vector.add(o);  
            }  
        }  
    }  
}
```

# Outline

- Motivation & examples
- Threads, shared memory, & synchronization
  - How do locks work?
- Data races (a lower-level property)
- How do data race detectors work?
- Atomicity (a higher-level property)
- **Concurrency exceptions & Summary**

Extra:

- Double-checked locking

# Concurrency exceptions?!

Java provides memory & type safety

- Buffer overflows, dangling pointers, array out-of-bounds, double frees, some memory leaks
- How are these handled? With exceptions?

# Concurrency exceptions?!

Java provides memory & type safety

- Buffer overflows, dangling pointers, array out-of-bounds, double frees, some memory leaks
- How are these handled? With exceptions?

Should languages (and the runtime systems & hardware that support them) provide **concurrency** correctness?

Check & enforce: atomicity, SC/DRF, determinism

# Summary

General-purpose parallel software: hard & unsolved

Challenging semantics for parallel programs

Understand how to write correct, scalable programs

→ only a few experts

# Outline

- Motivation & examples
- Threads, shared memory, & synchronization
  - How do locks work?
- Data races (a lower-level property)
- How do data race detectors work?
- Atomicity (a higher-level property)
- Concurrency exceptions & Summary

Extra:

- **Double-checked locking**



# Another example: double-checked locking

```
class Movie {  
    Vector<String> comments;  
  
    addComment(String s) {  
        if (comments == null) {  
            comments = new Vector<String>();  
        }  
        comments.add(s);  
    }  
}
```

# Another example: double-checked locking

```
class Movie {  
    Vector<String> comments;  
  
    addComment(String s) {  
        synchronized (this) {  
            if (comments == null) {  
                comments = new Vector<String>();  
            }  
        }  
        comments.add(s);  
    }  
}
```

# This is (incorrect) double-checked locking

```
class Movie {
    Vector<String> comments;

    addComment(String s) {
        if (comments == null) {
            synchronized (this) {
                if (comments == null) {
                    comments = new Vector<String>();
                }
            }
        }
        comments.add(s);
    }
}
```

# Another example: double-checked locking

```
addComment(String s) {  
    if (comments == null) {  
        synchronized (this) {  
            if (comments == null) {  
                comments =  
                new Vector<String>();  
            }  
        }  
    }  
}
```

```
addComment(String s) {
```

```
    if (comments == null) {  
        }  
        comments.add(s);  
    }
```

```
    comments.add(s);  
}
```

# Another example: double-checked locking


```
addComment(String s) {  
    if (comments == null) {  
        synchronized (this) {  
            if (comments == null) {  
                Vector temp =  
                alloc Vector;  
                temp.<init>();  
                comments = temp;  
            }  
        }  
    }  
}
```

```
addComment(String s) {  
  
    if (comments == null) {  
        }  
  
        comments.add(s);  
    }  
}
```

```
    comments.add(s);  
}
```

# Another example: double-checked locking

```
addComment(String s) {  
    if (comments == null) {  
        synchronized (this) {  
            if (comments == null) {  
                Vector temp =  
                alloc Vector;  
                temp.<init>();  
                comments = temp;  
            }  
        }  
    }  
}
```



```
addComment(String s) {  
  
    if (comments == null) {  
        }  
  
        comments.add(s);  
    }  
}
```

```
    comments.add(s);  
}
```

# Another example: double-checked locking

```
addComment(String s) {  
    if (comments == null) {  
        synchronized (this) {  
            if (comments == null) {  
                Vector temp =  
                alloc Vector;  
                comments = temp;  
  
                temp.<init>();  
            }  
        }  
    }  
    comments.add(s);  
}
```

```
addComment(String s) {  
  
  
  
  
  
  
  
  
    if (comments == null) {  
        }  
  
        comments.add(s);  
    }  
}
```

# This is correct double-checked locking

```
class Movie {  
    volatile Vector<String> comments;  
  
    addComment(String s) {  
        if (comments == null) {  
            synchronized (this) {  
                if (comments == null) {  
                    comments = new Vector<String>();  
                }  
            }  
        }  
        comments.add(s);  
    }  
}
```